


Deep Learning for Gophers



@iamrashminagpal



Rise of Deep Learning

**Google AI Tool
Identifies a
Tumor's
Mutations From
an Image**

**Pit.ai puts a financial twist on
reinforcement learning to
outperform hedge funds**

Identifying artificial intelligence “blind spots”


Finding a good read among billions of choices

As natural language processing techniques improve, suggestions are getting speedier and more relevant.

What is Deep Learning?


ARTIFICIAL INTELLIGENCE

Any technique that
enables computers to
mimic human behavior



MACHINE LEARNING

Ability to learn without
explicitly being
programmed



DEEP LEARNING

Extract patterns from
data using neural
networks


3 1 3 5 6 7
1 4 5 9 2 3

Why Deep Learning?

Hand engineered features are time consuming, brittle & not scalable in practise

Can we learn **underlying features** directly from data?

Low Level Features



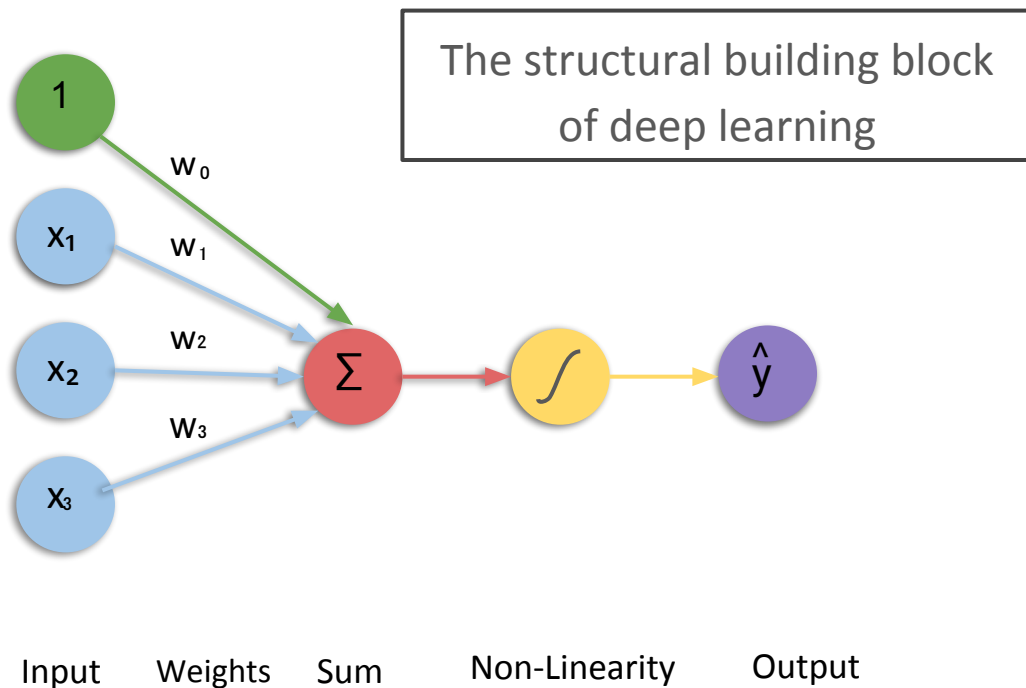
Mid Level Features



High Level Features



The Perceptron : Feedforward Propagation



Output

Linear combinations of input

$$\hat{y} = g \left(w_0 + \sum_{i=1}^m x_i w_i \right)$$

Non-linear activation function


Bias

```
package deep_learning_for_gophers

type Neuron struct {
    Act      ActivationFunc
    Input    []*Edge
    Output   []*Edge
    Value    float64
}

type Edge struct {
    Weight    float64
    Input     float64
    Output    float64
    IsBias    bool
}
```

The Perceptron : Feedforward Propagation



$$\hat{y} = g \left(w_0 + \sum_{i=1}^m x_i w_i \right)$$

$$\hat{y} = g (w_0 + \mathbf{X}^T \mathbf{W})$$

$$\text{where: } \mathbf{X} = \begin{bmatrix} x_1 \\ \vdots \\ x_m \end{bmatrix} \text{ and } \mathbf{W} = \begin{bmatrix} w_1 \\ \vdots \\ w_m \end{bmatrix}$$

```
package deep_learning_for_gophers
```


```
import ...
```

```
type initializeWeight func() float64
```

```
func NormalWeightInitialize(stdDev, mean float64) float64 {  
    return rand.NormFloat64()*stdDev + mean  
}
```

```
|  
func UniformWeightInitialize(stdDev, mean float64) float64 {  
    return (rand.Float64()-0.5)*stdDev + mean  
}
```


The Perceptron : Feedforward Propagation




Activation Function

$$\hat{y} = g(w_0 + X^T W)$$


- Example : Sigmoid function

$$g(z) = \sigma(z) = \frac{1}{1 + e^{-z}}$$



Common Activation Functions


Sigmoid Function



$$g(z) = \frac{1}{1 + e^{-z}}$$

$$g'(z) = g(z)(1 - g(z))$$


Hyperbolic Tangent



$$g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

$$g'(z) = 1 - g(z)^2$$

Rectified Linear Unit (ReLU)



$$g(z) = \max(0, z)$$

$$g'(z) = \begin{cases} 1, & z > 0 \\ 0, & \text{otherwise} \end{cases}$$

```
package deep_learning_for_gophers

import (
    "math"
    _ "math"
)

type ActivationFunc int

const(
    NoActivation ActivationFunc = 0
    SigmoidActivation ActivationFunc = 1
    TanhActivation ActivationFunc = 3
    ReLuActivation ActivationFunc = 4
    LinearActivation ActivationFunc = 5
    SoftMaxActivation ActivationFunc = 6
)

func GetActivationFunc(act ActivationFunc) Differentiable{
    switch act {
    case SigmoidActivation:
        return Sigmoid{}
    case ReLuActivation:
        return ReLU{}
    case LinearActivation:
```

```
func (s Sigmoid) Func(x float64) float64 {  
    return Logistic(x, a: 1)  
}
```

```
func (s Sigmoid) DFunc(y float64) float64{  
    return y*(1-y)  
}
```



```
type ReLU struct{}
```

```
func (a ReLU) Func(x float64) float64 { return math.Max(x, y: 0) }
```

```
func (a ReLU) DFunc(y float64) float64 {  
    if y > 0 : 1 ↗  
    return 0  
}
```


```
type linear struct{}
```

```
func (l linear) Func(x float64) float64 {return x}
```

```
func (l linear) DFunc(x float64) float64 {return 1}
```


Importance of Activation Functions

The purpose of activation functions is to **introduce non-linearities** into the network




What if we want to build a Neural Network to distinguish between blue vs orange points?

Importance of Activation Functions




Linear activations produce linear decisions
no matter what the network size



Non-linearities allow us to approximate
arbitrarily complex functions

The Perceptron : Example





We have: $w_0 = 1$ and $\mathbf{W} = \begin{bmatrix} 3 \\ -2 \end{bmatrix}$



$$\begin{aligned}\hat{y} &= g(w_0 + \mathbf{X}^T \mathbf{W}) \\ &= g\left(1 + \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}^T \begin{bmatrix} 3 \\ -2 \end{bmatrix}\right) \\ \hat{y} &= g(1 + 3x_1 - 2x_2)\end{aligned}$$

This is just a line in 2D!

The Perceptron : Example




The Perceptron : Example




Building Neural Networks with Perceptrons

The Perceptron : Simplified




The Perceptron : Simplified



$$z = w_0 + \sum_{j=1}^m x_j w_j$$

Single Layer Neural Network




Inputs

Hidden Layer

Final Output

$$z_i = w_{0,i}^{(1)} + \sum_{j=1}^m x_j w_{j,i}^{(1)} \quad \hat{y}_i = g \left(w_{0,i}^{(2)} + \sum_{j=1}^{d_1} z_j w_{j,i}^{(2)} \right)$$

Single Layer Neural Network



$$\begin{aligned} z_2 &= w_{0,2}^{(1)} + \sum_{j=1}^m x_j w_{j,2}^{(1)} \\ &= w_{0,2}^{(1)} + x_1 w_{1,2}^{(1)} + x_2 w_{2,2}^{(1)} + x_m w_{m,2}^{(1)} \end{aligned}$$

Applying Neural Networks

Example Problem


Will I be able to do justice to this topic?

Let's start with simple two feature model


X1 : Number of hours I practise for this talk

X2 : Number of hours spent on Netflix yesterday

Will I be able to do justice to this topic?




Single Layer Neural Network



Predicted : 0.1
Actual : 1

Quantifying Loss



Predicted : 0.1
Actual : 1

Training Neural Networks

Loss Optimization

We want to find the network weights that **achieve the lowest loss**

$$\mathbf{W}^* = \operatorname{argmin}_{\mathbf{W}} \frac{1}{n} \sum_{i=1}^n \mathcal{L}(f(x^{(i)}; \mathbf{W}), y^{(i)})$$


$$\mathbf{W}^* = \operatorname{argmin}_{\mathbf{W}} J(\mathbf{W})$$

Remember:

$$\mathbf{W} = \{\mathbf{W}^{(0)}, \mathbf{W}^{(1)}, \dots\}$$


Loss Optimization

$$W^* = \underset{w}{\operatorname{argmin}} J(w)$$




Loss Optimization

Randomly pick an initial (w_1, w_0)




Loss Optimization

Compute gradient $\frac{\partial J(w)}{\partial W}$




Loss Optimization

Take direction in opposite direction of gradient



Gradient Descent

Repeat until convergence



Gradient Descent

Algorithm

1. Initialize weights randomly


2. Loop until convergence

3. Compute gradient $\frac{\partial J(w)}{\partial W}$

4. Update weights $W \leftarrow W - \eta \frac{\partial J(w)}{\partial W}$


5. Return weights

Computing Gradients : Backpropagation



How does small change in one weight (eg w_2) affect the final loss $J(w)$?


Computing Gradients : Backpropagation



$$\frac{\partial J(W)}{\partial w_2}$$

Using chain rule


Computing Gradients : Backpropagation



$$\frac{\partial J(W)}{\partial w_2} = \frac{\partial J(W)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial w_2}$$

Using chain rule

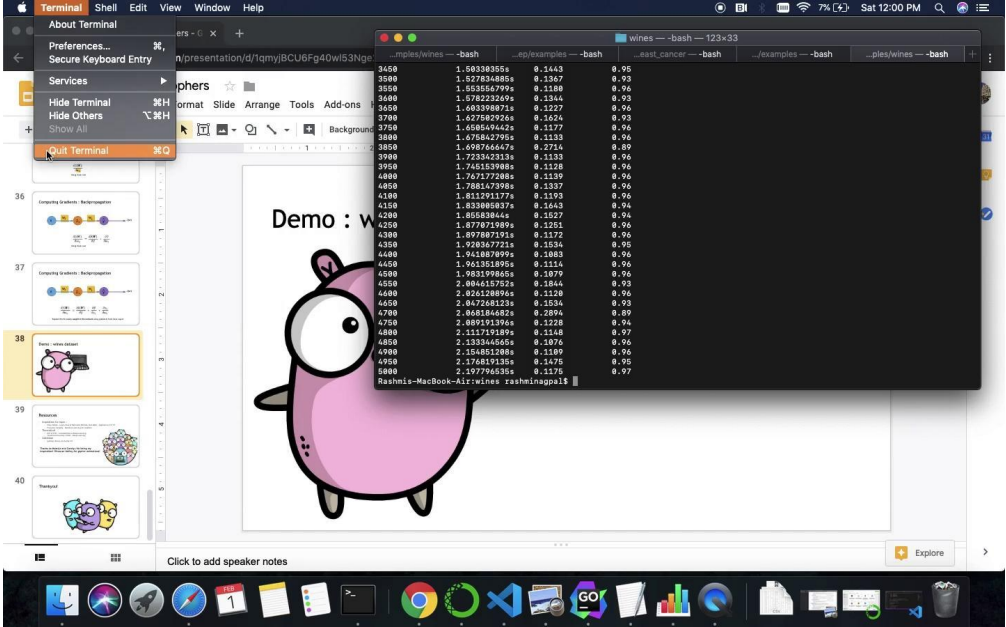

Computing Gradients : Backpropagation



$$\frac{\partial J(\mathbf{W})}{\partial w_1} = \frac{\partial J(\mathbf{W})}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1} * \frac{\partial z_1}{\partial w_1}$$

Repeat this for **every weight in the network** using gradients from later layers


Demo



Resources

- Inspiration for topic :
 - Ellen Kobes - Learn Neural Networks WithGo, Not Math : GopherCon EU'19
 - Francesc Campoy - Machine Learning for Gophers
- Theoretical
 - MIT 6.S191 - Introduction to Deep Learning
 - Stanford University CS230 - Deep Learning
- Codebase
 - godeep library as starter kit

Thanks to Maartje and Carolyn for being my inspiration! Ofcourse Ashley, for gopher animations!



Thankyou!

